

```
In [ ]: # 1.    Load the basic libraries and packages

import numpy as np
import tensorflow as tf
from tensorflow.keras.utils import to_categorical
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
```

```
In [ ]: # 2.    Load the dataset

from tensorflow.keras.datasets import fashion_mnist

# Load Fashion MNIST dataset
(train_X, train_Y), (test_X, test_Y) = fashion_mnist.load_data()

# Split data into training and validation sets
valid_X, valid_Y = train_X[:5000], train_Y[:5000]
train_X, train_Y = train_X[5000:], train_Y[5000:]
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
29515/29515 ————— 0s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
26421880/26421880 ————— 2s 0us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
5148/5148 ————— 0s 1us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
4422102/4422102 ————— 1s 0us/step

```
In [ ]: # 3.    Analyse the dataset

print(f"Training data shape: {train_X.shape}, Labels shape: {train_Y.shape}")
print(f"Validation data shape: {valid_X.shape}, Labels shape: {valid_Y.shape}")
print(f"Test data shape: {test_X.shape}, Labels shape: {test_Y.shape}")
print("Unique classes:", np.unique(train_Y))
```

Training data shape: (55000, 28, 28), Labels shape: (55000,)
Validation data shape: (5000, 28, 28), Labels shape: (5000,)
Test data shape: (10000, 28, 28), Labels shape: (10000,)
Unique classes: [0 1 2 3 4 5 6 7 8 9]

```
In [ ]: # 4.    Normalize the data

# Normalize pixel values to range 0-1
train_X = train_X.astype('float32') / 255.0
valid_X = valid_X.astype('float32') / 255.0
test_X = test_X.astype('float32') / 255.0
```

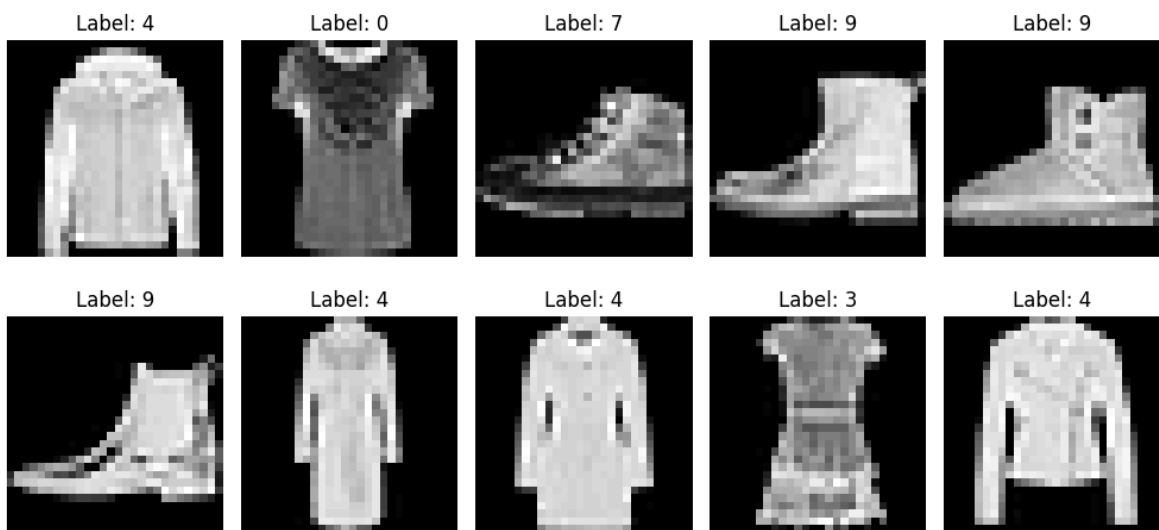
```
In [ ]: # 5.    Pre-process the data

# Reshape to include channel dimension
train_X = train_X.reshape(-1, 28, 28, 1)
valid_X = valid_X.reshape(-1, 28, 28, 1)
test_X = test_X.reshape(-1, 28, 28, 1)
```

```
# Convert labels to one-hot encoding
num_classes = 10
train_label = to_categorical(train_Y, num_classes)
valid_label = to_categorical(valid_Y, num_classes)
test_Y_one_hot = to_categorical(test_Y, num_classes)
```

In []: # 6. Visualize the Data

```
plt.figure(figsize=(10, 5))
for i in range(10):
    plt.subplot(2, 5, i + 1)
    plt.imshow(train_X[i].reshape(28, 28), cmap='gray')
    plt.title(f"Label: {train_Y[i]}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```



In []: # 7. Write the CNN model function

```
def create_cnn_model():
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Conv2D(32, (3, 3), activation='linear', padding='s
    model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
    model.add(tf.keras.layers.MaxPooling2D((2, 2), padding='same'))
    model.add(tf.keras.layers.Dropout(0.25))
    model.add(tf.keras.layers.Conv2D(64, (3, 3), activation='linear', padding='s
    model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
    model.add(tf.keras.layers.MaxPooling2D((2, 2), padding='same'))
    model.add(tf.keras.layers.Dropout(0.25))
    model.add(tf.keras.layers.Conv2D(128, (3, 3), activation='linear', padding='
    model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
    model.add(tf.keras.layers.MaxPooling2D((2, 2), padding='same'))
    model.add(tf.keras.layers.Dropout(0.4))
    model.add(tf.keras.layers.Flatten())
    model.add(tf.keras.layers.Dense(128, activation='linear'))
    model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
    model.add(tf.keras.layers.Dropout(0.3))
    model.add(tf.keras.layers.Dense(num_classes, activation='softmax'))
    return model
```

In []: # 8. Write the Cost Function

```
# Categorical Crossentropy as the Loss function
cost_function = tf.keras.losses.CategoricalCrossentropy()
```

In []: *# 9. Write the Gradient Descent optimization algorithm*

```
# Adam optimizer with default parameters
optimizer = tf.keras.optimizers.Adam()
```


In []: *# 10. Apply the training over the dataset to minimize the loss*


```
fashion_model = create_cnn_model()
fashion_model.compile(loss=cost_function, optimizer=optimizer, metrics=['accuracy'])


# Train the model
fashion_train = fashion_model.fit(
    train_X, train_label,
    batch_size=64, epochs=20,
    verbose=1, validation_data=(valid_X, valid_label)
)
```


```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
```


```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
/usr/local/lib/python3.10/dist-packages/keras/src/layers/activations/leaky_relu.py:41: UserWarning: Argument `alpha` is deprecated. Use `negative_slope` instead.
    warnings.warn(
```


Epoch 1/20
860/860  **118s** 132ms/step - accuracy: 0.6928 - loss: 0.8301 - val_accuracy: 0.8812 - val_loss: 0.3291


Epoch 2/20
860/860  **142s** 133ms/step - accuracy: 0.8584 - loss: 0.3817 - val_accuracy: 0.9016 - val_loss: 0.2733


Epoch 3/20
860/860  **141s** 132ms/step - accuracy: 0.8792 - loss: 0.3272 - val_accuracy: 0.8968 - val_loss: 0.2619


Epoch 4/20
860/860  **144s** 135ms/step - accuracy: 0.8881 - loss: 0.3022 - val_accuracy: 0.9152 - val_loss: 0.2320


Epoch 5/20
860/860  **114s** 133ms/step - accuracy: 0.8979 - loss: 0.2751 - val_accuracy: 0.9186 - val_loss: 0.2210


Epoch 6/20
860/860  **111s** 129ms/step - accuracy: 0.9000 - loss: 0.2611 - val_accuracy: 0.9202 - val_loss: 0.2164


Epoch 7/20
860/860  **141s** 127ms/step - accuracy: 0.9068 - loss: 0.2490 - val_accuracy: 0.9168 - val_loss: 0.2238


Epoch 8/20
860/860  **111s** 129ms/step - accuracy: 0.9103 - loss: 0.2398 - val_accuracy: 0.9238 - val_loss: 0.2036


Epoch 9/20
860/860  **143s** 130ms/step - accuracy: 0.9096 - loss: 0.2374 - val_accuracy: 0.9202 - val_loss: 0.2103


Epoch 10/20
860/860  **114s** 133ms/step - accuracy: 0.9129 - loss: 0.2298 - val_accuracy: 0.9258 - val_loss: 0.1988


Epoch 11/20
860/860  **115s** 133ms/step - accuracy: 0.9162 - loss: 0.2227 - val_accuracy: 0.9238 - val_loss: 0.2044


Epoch 12/20
860/860  **142s** 134ms/step - accuracy: 0.9155 - loss: 0.2247 - val_accuracy: 0.9284 - val_loss: 0.1954


Epoch 13/20
860/860  **114s** 132ms/step - accuracy: 0.9178 - loss: 0.2180 - val_accuracy: 0.9232 - val_loss: 0.2054


Epoch 14/20
860/860  **141s** 131ms/step - accuracy: 0.9196 - loss: 0.2125 - val_accuracy: 0.9260 - val_loss: 0.1952


Epoch 15/20
860/860  **117s** 136ms/step - accuracy: 0.9184 - loss: 0.2105 - val_accuracy: 0.9246 - val_loss: 0.2030

Epoch 16/20
860/860  **114s** 132ms/step - accuracy: 0.9215 - loss: 0.2067 - val_accuracy: 0.9314 - val_loss: 0.1901

Epoch 17/20
860/860  **117s** 136ms/step - accuracy: 0.9245 - loss: 0.2027 - val_accuracy: 0.9300 - val_loss: 0.1901

Epoch 18/20
860/860  **116s** 135ms/step - accuracy: 0.9248 - loss: 0.1973 - val_accuracy: 0.9328 - val_loss: 0.1891

Epoch 19/20
860/860  **147s** 140ms/step - accuracy: 0.9277 - loss: 0.1933 - val_accuracy: 0.9304 - val_loss: 0.1850


Epoch 20/20
860/860  **139s** 136ms/step - accuracy: 0.9268 - loss: 0.1937 - val_accuracy: 0.9312 - val_loss: 0.1869


```
In [ ]: # 11. Observe the cost function vs iterations learning curve
```

```
accuracy = fashion_train.history['accuracy']
val_accuracy = fashion_train.history['val_accuracy']
loss = fashion_train.history['loss']
val_loss = fashion_train.history['val_loss']
epochs = range(len(accuracy))

# Accuracy Curve
plt.plot(epochs, accuracy, 'bo', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.figure()

# Loss Curve
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.show()
```





Result

In []: *# a. Training dataset*

```
print("Training data shape:", train_X.shape)
print("Validation data shape:", valid_X.shape)
print("Test data shape:", test_X.shape)
print("Number of classes:", num_classes)
print("Unique classes:", np.unique(train_Y))
```

```
Training data shape: (55000, 28, 28, 1)
Validation data shape: (5000, 28, 28, 1)
Test data shape: (10000, 28, 28, 1)
Number of classes: 10
Unique classes: [0 1 2 3 4 5 6 7 8 9]
```


In []: *# b. Model summary*

```
# Before Regularization
model = tf.keras.Sequential()
model.add(tf.keras.layers.Conv2D(32, (3, 3), activation='linear', padding='same'))
model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
model.add(tf.keras.layers.MaxPooling2D((2, 2), padding='same'))
model.add(tf.keras.layers.Conv2D(64, (3, 3), activation='linear', padding='same'))
model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
model.add(tf.keras.layers.MaxPooling2D((2, 2), padding='same'))
model.add(tf.keras.layers.Conv2D(128, (3, 3), activation='linear', padding='same'))
model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
model.add(tf.keras.layers.MaxPooling2D((2, 2), padding='same'))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation='linear'))
model.add(tf.keras.layers.LeakyReLU(alpha=0.1))
```

```
model.add(tf.keras.layers.Dense(num_classes, activation='softmax'))
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	
conv2d_3 (Conv2D)	(None, 28, 28, 32)	
leaky_re_lu_4 (LeakyReLU)	(None, 28, 28, 32)	
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 32)	
conv2d_4 (Conv2D)	(None, 14, 14, 64)	
leaky_re_lu_5 (LeakyReLU)	(None, 14, 14, 64)	
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 64)	
conv2d_5 (Conv2D)	(None, 7, 7, 128)	
leaky_re_lu_6 (LeakyReLU)	(None, 7, 7, 128)	
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 128)	
flatten_1 (Flatten)	(None, 2048)	
dense_2 (Dense)	(None, 128)	
leaky_re_lu_7 (LeakyReLU)	(None, 128)	
dense_3 (Dense)	(None, 10)	



Total params: 356,234 (1.36 MB)

Trainable params: 356,234 (1.36 MB)


Non-trainable params: 0 (0.00 B)


```
In [ ]: # c. Training and validation accuracy w.r.t epochs before regularization


model.compile(loss=tf.keras.losses.CategoricalCrossentropy(), optimizer=tf.keras


# Train model without regularization
train_no_reg = model.fit(train_X, train_label, batch_size=64, epochs=20, validat


plt.plot(train_no_reg.history['accuracy'], label='Training Accuracy')
plt.plot(train_no_reg.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Before Regularization')
plt.legend()
plt.show()
```


Epoch 1/20
860/860  **111s** 126ms/step - accuracy: 0.7685 - loss: 0.6339 - val_accuracy: 0.8868 - val_loss: 0.3094


Epoch 2/20
860/860  **142s** 126ms/step - accuracy: 0.8993 - loss: 0.2801 - val_accuracy: 0.9076 - val_loss: 0.2627


Epoch 3/20
860/860  **144s** 128ms/step - accuracy: 0.9132 - loss: 0.2340 - val_accuracy: 0.9136 - val_loss: 0.2332


Epoch 4/20
860/860  **142s** 128ms/step - accuracy: 0.9227 - loss: 0.2067 - val_accuracy: 0.9166 - val_loss: 0.2294


Epoch 5/20
860/860  **140s** 125ms/step - accuracy: 0.9337 - loss: 0.1766 - val_accuracy: 0.9190 - val_loss: 0.2171


Epoch 6/20
860/860  **138s** 121ms/step - accuracy: 0.9457 - loss: 0.1498 - val_accuracy: 0.9254 - val_loss: 0.2103


Epoch 7/20
860/860  **145s** 125ms/step - accuracy: 0.9513 - loss: 0.1305 - val_accuracy: 0.9192 - val_loss: 0.2341


Epoch 8/20
860/860  **139s** 122ms/step - accuracy: 0.9612 - loss: 0.1060 - val_accuracy: 0.9158 - val_loss: 0.2474


Epoch 9/20
860/860  **145s** 125ms/step - accuracy: 0.9672 - loss: 0.0900 - val_accuracy: 0.9202 - val_loss: 0.2529


Epoch 10/20
860/860  **142s** 125ms/step - accuracy: 0.9704 - loss: 0.0779 - val_accuracy: 0.9240 - val_loss: 0.2696


Epoch 11/20
860/860  **143s** 126ms/step - accuracy: 0.9752 - loss: 0.0656 - val_accuracy: 0.9212 - val_loss: 0.2907


Epoch 12/20
860/860  **138s** 121ms/step - accuracy: 0.9791 - loss: 0.0564 - val_accuracy: 0.9238 - val_loss: 0.3015


Epoch 13/20
860/860  **107s** 124ms/step - accuracy: 0.9806 - loss: 0.0491 - val_accuracy: 0.9164 - val_loss: 0.3317


Epoch 14/20
860/860  **143s** 125ms/step - accuracy: 0.9845 - loss: 0.0406 - val_accuracy: 0.9204 - val_loss: 0.3473


Epoch 15/20
860/860  **109s** 127ms/step - accuracy: 0.9864 - loss: 0.0367 - val_accuracy: 0.9224 - val_loss: 0.3637


Epoch 16/20
860/860  **142s** 128ms/step - accuracy: 0.9874 - loss: 0.0353 - val_accuracy: 0.9242 - val_loss: 0.3626

Epoch 17/20
860/860  **115s** 134ms/step - accuracy: 0.9858 - loss: 0.0376 - val_accuracy: 0.9250 - val_loss: 0.3952

Epoch 18/20
860/860  **136s** 127ms/step - accuracy: 0.9886 - loss: 0.0301 - val_accuracy: 0.9210 - val_loss: 0.4122


Epoch 19/20
860/860  **144s** 130ms/step - accuracy: 0.9897 - loss: 0.0283 - val_accuracy: 0.9122 - val_loss: 0.4835

Epoch 20/20
860/860  **135s** 122ms/step - accuracy: 0.9913 - loss: 0.0251 - val_accuracy: 0.9226 - val_loss: 0.4272



```
In [ ]: # d.    Training and validation loss w.r.t epochs before regularization

plt.plot(train_no_reg.history['loss'], label='Training Loss')
plt.plot(train_no_reg.history['val_loss'], label='Validation Loss')
plt.title('Loss Before Regularization')
plt.legend()
plt.show()
```





```
In [ ]: # e.    Training and validation accuracy w.r.t epochs after regularization


# Model with Dropout Regularization
reg_model = create_cnn_model()
reg_model.compile(loss=tf.keras.losses.CategoricalCrossentropy(), optimizer=tf.k


train_with_reg = reg_model.fit(train_X, train_label, batch_size=64, epochs=20, v


plt.plot(train_with_reg.history['accuracy'], label='Training Accuracy')
plt.plot(train_with_reg.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy After Regularization')
plt.legend()
plt.show()
```


Epoch 1/20
860/860  **121s** 138ms/step - accuracy: 0.6963 - loss: 0.8150 - val_accuracy: 0.8718 - val_loss: 0.3397


Epoch 2/20
860/860  **142s** 138ms/step - accuracy: 0.8599 - loss: 0.3797 - val_accuracy: 0.8930 - val_loss: 0.2973


Epoch 3/20
860/860  **143s** 139ms/step - accuracy: 0.8768 - loss: 0.3310 - val_accuracy: 0.9014 - val_loss: 0.2579


Epoch 4/20
860/860  **117s** 136ms/step - accuracy: 0.8881 - loss: 0.2994 - val_accuracy: 0.9054 - val_loss: 0.2487


Epoch 5/20
860/860  **144s** 138ms/step - accuracy: 0.8967 - loss: 0.2809 - val_accuracy: 0.9134 - val_loss: 0.2255


Epoch 6/20
860/860  **139s** 135ms/step - accuracy: 0.9000 - loss: 0.2628 - val_accuracy: 0.9146 - val_loss: 0.2390


Epoch 7/20
860/860  **145s** 138ms/step - accuracy: 0.9048 - loss: 0.2553 - val_accuracy: 0.9176 - val_loss: 0.2216


Epoch 8/20
860/860  **117s** 136ms/step - accuracy: 0.9091 - loss: 0.2447 - val_accuracy: 0.9222 - val_loss: 0.2044


Epoch 9/20
860/860  **118s** 137ms/step - accuracy: 0.9117 - loss: 0.2387 - val_accuracy: 0.9270 - val_loss: 0.2010


Epoch 10/20
860/860  **141s** 136ms/step - accuracy: 0.9124 - loss: 0.2332 - val_accuracy: 0.9198 - val_loss: 0.2068


Epoch 11/20
860/860  **145s** 139ms/step - accuracy: 0.9174 - loss: 0.2211 - val_accuracy: 0.9270 - val_loss: 0.1953


Epoch 12/20
860/860  **137s** 134ms/step - accuracy: 0.9174 - loss: 0.2222 - val_accuracy: 0.9198 - val_loss: 0.2108


Epoch 13/20
860/860  **119s** 139ms/step - accuracy: 0.9183 - loss: 0.2162 - val_accuracy: 0.9252 - val_loss: 0.2093


Epoch 14/20
860/860  **116s** 135ms/step - accuracy: 0.9204 - loss: 0.2102 - val_accuracy: 0.9250 - val_loss: 0.2030


Epoch 15/20
860/860  **148s** 142ms/step - accuracy: 0.9189 - loss: 0.2103 - val_accuracy: 0.9256 - val_loss: 0.1996


Epoch 16/20
860/860  **117s** 136ms/step - accuracy: 0.9243 - loss: 0.2052 - val_accuracy: 0.9302 - val_loss: 0.1886

Epoch 17/20
860/860  **143s** 138ms/step - accuracy: 0.9244 - loss: 0.2006 - val_accuracy: 0.9280 - val_loss: 0.1997

Epoch 18/20
860/860  **123s** 143ms/step - accuracy: 0.9258 - loss: 0.1988 - val_accuracy: 0.9290 - val_loss: 0.1983


Epoch 19/20
860/860  **136s** 136ms/step - accuracy: 0.9241 - loss: 0.1985 - val_accuracy: 0.9238 - val_loss: 0.2054

Epoch 20/20
860/860  **121s** 141ms/step - accuracy: 0.9245 - loss: 0.1991 - val_accuracy: 0.9232 - val_loss: 0.1984



```
In [ ]: # f.    Training and validation loss w.r.t epochs after regularization

plt.plot(train_with_reg.history['loss'], label='Training Loss')
plt.plot(train_with_reg.history['val_loss'], label='Validation Loss')
plt.title('Loss After Regularization')
plt.legend()
plt.show()
```



```
In [ ]: # g.    Original v/s predicted labels for correct predicted observations

predicted_classes = np.argmax(reg_model.predict(test_X), axis=1)
correct_indices = np.where(predicted_classes == test_Y[0])

print(f"Number of Correct Predictions: {len(correct_indices)}")
for i, correct in enumerate(correct_indices[:9]):
    plt.subplot(3, 3, i + 1)
    plt.imshow(test_X[correct].reshape(28, 28), cmap='gray')
    plt.title(f"Pred: {predicted_classes[correct]}, True: {test_Y[correct]}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```

313/313 ————— 6s 19ms/step
 Number of Correct Predictions: 9187


Pred: 9, True: 9




Pred: 2, True: 2



Pred: 1, True: 1




Pred: 1, True: 1



Pred: 6, True: 6




Pred: 1, True: 1




Pred: 4, True: 4



Pred: 6, True: 6



Pred: 5, True: 5




```
In [ ]: # h.    Original v/s predicted labels for incorrect predicted observations

incorrect_indices = np.where(predicted_classes != test_Y)[0]

print(f"Number of Incorrect Predictions: {len(incorrect_indices)}")
for i, incorrect in enumerate(incorrect_indices[:9]):
    plt.subplot(3, 3, i + 1)
    plt.imshow(test_X[incorrect].reshape(28, 28), cmap='gray')
    plt.title(f"Pred: {predicted_classes[incorrect]}, True: {test_Y[incorrect]}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```

Number of Incorrect Predictions: 813

Pred: 5, True: 7



Pred: 6, True: 4



Pred: 5, True: 9




Pred: 2, True: 4



Pred: 0, True: 6



Pred: 6, True: 3



Pred: 6, True: 2




Pred: 6, True: 4



Pred: 3, True: 2



 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To understand the process of convolution over the image and apply over the classification problem	
Experiment No: 4	Date:	Enrolment No: 92200133003

Aim: To understand the process of convolution over the image and apply over the classification problem


IDE: Google Colab

Theory:

Convolutional Neural Networks (CNN) are complex feed forward neural networks. CNNs are used for image classification and recognition because of its high accuracy. There are three types of layers in a convolutional neural network: i. Convolutional layer ii. Pooling layer iii. Fully connected layer. Each of these layers has different parameters that can be optimized and performs a different task on the input data.


What is Pooling Layer?

Pooling layer is responsible for reducing the spatial size of the Convolved Feature. This is to decrease the computational power required to process the data through dimensionality reduction. There are two types of Pooling: i. Average Pooling. ii. Max Pooling. Max Pooling returns the maximum value from the portion of the image covered by the Kernel. On the other hand, Average Pooling returns the average of all the values from the portion of the image covered by the Kernel. Max Pooling also performs as a Noise Suppressant. It discards the noisy activations altogether and also performs de-noising along with dimensionality reduction. On the other hand, Average Pooling simply performs dimensionality reduction as a noise suppressing mechanism. Hence, we can say that Max Pooling performs a lot better than Average Pooling.



What is Convolutional Layer?

Convolutional layers are the major building blocks used in convolutional neural networks. A convolution is the simple application of a filter to an input that results in an activation. Repeated application of the same filter to an input results in a map of activations called a feature map, indicating the locations and strength of a detected feature in an input, such as an image. A convolutional layer contains a set of filters whose parameters need to be learned. The height and weight of the filters are smaller than those of the input volume. Each filter is convolved with the input volume to compute an activation map made of neurons.

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To understand the process of convolution over the image and apply over the classification problem	
Experiment No: 4	Date:	Enrolment No: 92200133003


What is Fully Connected Layer?

A fully connected layer that takes the output of convolution/pooling and predicts the best label to describe the image. We have three layers in the full connection step: i. Input layer, ii. Fully-connected layer, iii. Output layer.

Input Layer: It takes the output of the previous layers, “flattens” them and turns them into a single vector that can be an input for the next stage.

Fully Connected Layer: It takes the inputs from the feature analysis and applies weights to predict the correct label.


Output Layer: It gives the final probabilities for each label.



ReLU Layer: ReLU is an activation function. Rectified Linear Unit (ReLU) transform function only activates a node if the input is above a certain quantity, while the input is below zero, the output is zero, but when the input rises above a certain threshold, it has a linear relationship with the dependent variable. The main aim is to remove all the negative values from the convolution. All the positive values remain the same but all the negative values get changed to zero.

Methodology:

1. Load the basic libraries and packages
2. Load the dataset
3. Analyse the dataset
4. Normalize the data
5. Pre-process the data
6. Visualize the Data

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To understand the process of convolution over the image and apply over the classification problem	
Experiment No: 4	Date:	Enrolment No: 92200133003

7. Write the CNN model function
8. Write the Cost Function
9. Write the Gradient Descent optimization algorithm
10. Apply the training over the dataset to minimize the loss
11. Observe the cost function vs iterations learning curve

Program (Code):

To be attached with

Results:

To be attached with


- a. Training dataset
- b. Model summary
- c. Training and validation accuracy w.r.t epochs before regularization
- d. Training and validation loss w.r.t epochs before regularization
- e. Training and validation accuracy w.r.t epochs after regularization
- f. Training and validation loss w.r.t epochs after regularization
- g. Original v/s predicted labels for correct predicted observations
- h. Original v/s predicted labels for incorrect predicted observations

Pre Lab Exercise:

- a. What is Convolution process? Explain giving example.

- b. What are different layers in CNN model?

- c. What is the requirement of the pooling layer in the CNN model?

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To understand the process of convolution over the image and apply over the classification problem	
Experiment No: 4	Date:	Enrolment No: 92200133003

d. What is the requirement of the use of ReLU activation function after convolution step?


Observation and Result Analysis:

e. Nature of the dataset

f. Training Process without regularization

g. After regularization in the training Process

h. Observation over the Learning Curves

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To understand the process of convolution over the image and apply over the classification problem	
Experiment No: 4	Date:	Enrolment No: 92200133003

Post Lab Exercise:

- a. Why CNN is preferred over ANN for images

- b. Can CNN be applied over Text data? If yes, then how. If no, then why?

- c. What is the role of dropout layer?

- d. What will happen if maxpooling is replaced with minpooling?
